

COMP2010 Compilers Lexing and Parsing Coursework

Submission Deadline

Friday 27th February 2015 @ 11:55PM

The goal of this 2010 Compiler coursework is to build a lexer and parser for the S++ programming language. Use JFlex 1.6.0 and Cup version 11b-20141204, using *only* the specified versions, to automatically generate code for your scanner and parser¹. This coursework broadly requires writing regular expressions covering all legal words in the language (`Lexer.lex`), and a context free grammar describing its rules (`Parser.cup`).

A single mark will be given to this coursework, comprising 9% of your mark for the COMP2010 module. Please Submit your work for (JFlex/CUP specifications) on Friday 27th February 2015 @ 11:55PM.

Detailed submission instructions are given at the end of the document.

1 The S++ Language

You are to build a lexer and a parser for S++.

A program in S++ consists of a list declarations, one of which is a **main** function. The declaration list defines global variables and new data types as well as functions, but cannot be empty: it must have a **main**.

S++ has two types of **comments**. First, any text, other than a newline, following `#` to the end of the line is a comment. Second, any text, including newlines, enclosed within `/# ... #/` is a comment, and may span multiple lines.

An **identifier** starts with a letter, followed by an arbitrary number of underscores, letters, or digits. Identifiers are case-sensitive. Punctuation other than underscore is not allowed.

A **character** is a single letter, punctuation symbol, or digit wrapped in `' '` and has type **char**.

The **boolean constants** are **T** and **F** and have type **bool**.

Numbers are integers (type **int**), rationals (type **rat**), or floats (type **float**). Negative numbers are represented by the `' - '` symbol before the digits. Examples of integers include `1` and `-1234`; Examples of rationals include `1/3` and `-345_11/3`; examples of floats are `-0.1` and `3.14`.

A **Dictionary** (type **dict**) is a collection of (**key**, **value**) pairs, with the constraint that a key appears at most once in the collection. When declaring a dictionary, one must specify the type of the keys and values. For example, `d : dict<int, char> := {key1:val1, key2:val2, ...}`; here, the keys must be integers and the values, characters. Use the special type keyword **top** to define a dictionary that allows any type for a key or value: `d : dict<int, top> := {1:1, 2:'c', 7:3/5, {1:T}}`. An empty dictionary is `{}`. The assignment `d[k] := v` binds `k` to `v` in `d`. If `d` already contains `k`, `k` is rebound to `v`; if not, the pair (`k`, `v`) is added to `d` and accessed by `d[k]`.

Sequences (type **seq**) are ordered containers of elements. Sequences have nonnegative length. A sequence has a type: its declaration specifies the type of elements it contains. For instance, `l : seq<int> := [1,2,3]`, and `str : seq<char> := ['f', 'r', 'e', 'd', 'd', 'y']`. As with **dict** above, you can use the **top** keyword to specify a sequence that contains any type, writing `s : seq<top> := [1, 1/2, 3.14, ['f', 'o', 'u', 'r']]`; The zero length list is `[]`.

S++ sequences support the standard **indexing** syntax. For any sequence `s`, the operator `len(s) : seq → ℕ` returns the length of `s` and the indices of `s` range from 0 to `len(s) - 1`. The expression `s[index]` returns the element in `s` at `index`. String literals are syntactic sugar for character sequences, so `"abc"` is `['a', 'b', 'c']`. For the sequence `s : seq<char> := "hello world"`, `s[len(s)-1] := 'd'`, where `len(s) = 11`.

¹Section 3 explains why I have imposed these constraints on the permitted versions of these tools.

Primitive Data Types **bool, int, rat, float, char**
Aggregate Data Types **dict, seq**

Table 1: S++ data types.

Operator	Defined Over	Syntax
Boolean	bool	!, &&, , ==>
Numeric	int, rat, float	+ - * / ^
Dictionary	dict	in , d[k], len(d)
Sequence	seq	in , ::, len(s), s[i], s[i:j], s[i:], s[:i]
Comparison	Numeric	< <=
	Boolean, Numeric	= !=

Table 2: S++ operators.

Sequences in S++ also support **sequence slicing** as in languages like Python or Ruby: `id[i:j]` returns another sequence, which is a subsequence of `id` starting at `id[i]` and ending at `id[j]`. Given `a = [1,2,3,4,5]`, `a[1:3]` is `[2,3,4]`. When the start index is not given, it implies that the subsequence starts from index 0 of the original sequence (e.g., `a[:2]` is `[1,2]`). Similarly, when the end index is not given, the subsequence ends with the last element of the original sequence (e.g., `a[3:]` is `[4,5]`). Finally, indices can be negative, in which case its value is determined by counting backwards from the end of the original sequence: `a[2:-1]` is equivalent to `a[2:len(a)-1]` and, therefore, is `[3,4,5]`, while `s[-2]` is 4. The lower index in a slice must be positive and smaller than the upper index, after the upper index has been subtracted from the sequences length if it was negative.

Table 1 defines S++’s builtin data types. In Table 2, “!” denotes logical not, “&&” logical and and “||” logical or, as is typical in the C language family; “==>” denotes implication. Note that “=” is referential equality, *not* the assignment operator in S++. The **in** operator checks whether an element (key) is present in a sequence (dictionary), as in `2 in [1,2,3]` or `2 in {1:"one", 2:"two"}`, and returns a boolean. Note that **in** only operates on the outermost sequence: `3 in [[1],[2],[3]]` is F, or false. “::” operator denotes concatenation, “`s[i]`” returns the *i*th entry in `s` and “`len(s)`” returns the length of `s` as defined in the discussion of sequences and their indexing above.

1.1 Declarations

The syntax of field or variable declaration is “`id : type`”. A data type declaration is

```
tdef type_id { declaration_list } ;
```

where `declaration_list` is a colon-separated list of field/variable declarations. Once declared, a data type can be used as a type in subsequent declarations. For readability, S++ supports type aliasing: the directive “`alias old_name new_name ;`” can appear in a declaration list and allows the use of `new_name` in place of `old_name`.

```
alias seq<char> string;
tdef person { name:string, surname:string, age:int };
tdef family { mother:person, father:person, children:seq<person> };
```

Listing 1: S++ data type declaration examples.

Function In Listing 2, the function’s **name** is an identifier. Each formal parameter follows the variable/field declaration syntax, `id : type`. The formal parameter list is comma-separated list of parameter declarations. A function’s body consists of local variable declarations, if any, followed by statements. The return type of the function is `returnType` and is omitted when the function does not return a value.

```
fdef name (formal_parameter_list) { body } : returnType ;
fdef name (formal_parameter_list) { body } ;
```

Listing 2: S++ function declaration syntax.

p.age + 10	Assumes “p: person;” previously declared
b - foo(sum(10, c), bar()) = 30	Illustrates method calls
s1 :: s2 :: [1,2]	Assumes s1 and s2 have type seq<int>

Table 3: S++ expression examples.

Expressions

S++ expressions are applications of the operators defined above. Parentheses enforce precedence. For user-defined data type definitions, field references are expressions and have the form `id.field`. Function calls can be either a statement or an expression; their parameters are expressions that, in the semantic phase (*i.e.* not this coursework), would be required to produce a type that can unify with the type of their parameter. Table 3 contains example expressions.

Statements

In Table 4, **var** indicates a variable. An **expression_list** is a comma-separated list of expressions. As above, a body consists of local variable declarations (if any), followed by statements. Statements, apart from **if-else** and **while**, terminate with a semicolon. The return statement appears in a function body, where it is optional.

Assignment	var := expression ;
Input	read var ;
Output	print expression ;
Function Call	functionId (expression_list) ;
	if (expression) then body fi
	if (expression) then body else body fi
Control Flow	while (expression) do body od
	forall (item in iterable) do body od
	return expression ;

Table 4: S++ statements.

Variables may be initialised at the time of declaration: “`id : type := init ;`”. For newly defined data types, initialisation consists of a sequence of comma-separated values, each of which is assigned to the data type fields in the order of declaration. Listing 3 contains examples.

The statement **read** `var`; reads a value from the standard input and stores it in `var`; the statement **print** prints evaluation of its expression parameter, followed by a newline. The **if** and **while** statements behave like those in the C family language. The iterable in **forall** is either a sequence or a dictionary; `item` is bound to each element in order for a sequence and in an arbitrary order for a dictionary. Listing 4 shows how to use **forall**.

Listing 5 shows an example function. The function **main** is the special S++ function where execution starts. S++’s **main** returns no value. Listing 6 contains an example.

```

a : dict<int, char> := { 1:'1', 2:'2', 3:'3' } ;
b : int := 10;
c : string := "hello world!";
d : person := "Shin", "Yoo", 30;
e : char := 'a';
f : seq<rat> := [ 1/2, 3, 4_2/17, -7 ];

```

Listing 3: S++ variable declaration and initialization examples.

```

a : seq<int> := [1, 2, 3];
forall(n in a) do
  print n * 2;
od

```

Listing 4: S++ iterables example.

```

fdef reverse (inseq : seq<top>) {
  outseq : seq<top> := [];
  i : int := 0;
  while (i < len(l)) do
    outseq := inseq[i] :: outseq;
    i := i + 1;
  od
  return outseq;
} : seq<top> ;

```

Listing 5: S++ example function.

```

main {
  a : seq<int> := [1,2,3];
  b : seq<int> := reverse(a);
}

```

Listing 6: S++ main section.

2 Error Handling

Your parser will be tested against a test suite of positive and negative tests. This testing is scripted; so it is important for your output to match what the script expects. Add the following function definition into the "parser code" section of your Cup file, between its `{ : nd : }` delimiters.

```
public void syntax_error(Symbol current_token) {
    report_error(
        "Syntax error at line " + (current_token.left+1) + ", column "
        + current_token.right, null
    );
}
```

Listing 7: S++ compiler error message format.

3 Submission Requirements and Instructions

Your scanner (lexer) must

- Use `JLex` (or `JFlex`) to automatically generate a scanner for the S++ language;
- Make use of macro definitions where necessary. Choose meaningful token type names to make your specification readable and understandable;
- Ignore whitespace and comments; and
- Report the line and the column (offset into the line) where an error, usually unexpected input, first occurred. Use the code in Section 2, which specifies the format that will be matched by the grading script.

Your parser must

- Use `CUP` to automatically produce a parser for the S++ language;
- Resolve ambiguities in expressions using the precedence and associativity rules;
- Print "parsing successful", followed by a newline, if the program is syntactically correct.

Your scanner and parser must work together.

Once the scanner and parser have been successfully produced using `JFlex` and `CUP`, write a `QC.java` class similar to the `Test.java` seen during lecture, to test your code on the test files provided on the course webpage.

I have provided a makefile on Moodle. This makefile *must* build your project, from source, when `make` is issued,

- on Ubuntu 14.04.01 64bit
- using `JFlex` 1.6.0
- using `Cup` version 11b-20141204
- using Java SE 8 Update 31

If your submission fails to build using this Makefile on Ubuntu 14.04.01 64bit with these versions of the tools, your mark will be zero.

The provided makefile has a test rule. The marking script will call this rule to run your parser against a set of test cases. This set of test cases includes public test cases provided via Moodle and private ones; they include positive tests, on which your parser must emit "parsing successful" followed by a newline and *nothing else*, and negative tests on which your parser must emit the correct line and column of the error, as specified in Section 2 above.

Your mark will be the number of positive tests cases you correctly pass and the number of negative test cases you correctly fail divided by the total number of test cases.

Via Moodle, each group should submit a tar ball, or zip file, that contains

- The Makefile with which I have provided you
- The JFlex specification `Lexer.lex`
- The CUP specification `Parser.cup`
- Any other classes you have defined, if any, using the directory layout the Makefile expects
- To save space, it is not necessary zip up and include the JFlex and Cup jars in the lib directory

Only one submission per group is necessary.

The deadline for completion of this part of the coursework is Friday 27th February 2015 @ 11:55PM. Any coursework handed in later than 2 working days after the deadline will automatically receive a zero mark.